

An Assignment Report

On

Review of Application Execution in Python: Case Study

by

Aditya Sharma

AU18B1009

Under the guidance of

Satyajit Pangaonkar



School of Engineering

Avantika University, Ujjain

2020-2021

Content:	Page
1. Introduction	2
2. Title	3
3. Statement	3
4. Identification of dataset	3 - 15
5. Identification of Learning Model (Supervised Learning)	16 - 18
6. Key Learning and outcomes	19

Introduction:

The used car market has significantly grown in recent times, with clients ranging from used car dealers and buyers. Used cars might not be as good as new ones, but when one, running short on a budget, it can be the best option for a while. The Indian used car market is segmented by the organized and unorganized segments. However, C2C (customer to customer) channel has also been used for the sales of pre-owned cars in the market. Also, some service providers are commercially available to guide and providing a platform for the user. There are many factors that involved attention while purchasing the used car.

Assignment No-03 (CS 6301- Machine Learning)

Title: Used Car condition prediction

Statement:

To create a robust model that allows stakeholders to predict the condition of a used vehicle.

A. Identification of the Dataset :

I. Type of the Dataset (Description) : Multivariate, Structured dataset

The dataset is related to kaggle compete dataset. It's a limited participant competition dataset. There are two datasets associated with the competition, train-data.csv and test-data.csv.

Train-data.csv: (Rows: 5975 , Columns: 17)

Test-data.csv: (Rows: 1201, Columns: 12)

II. Data Quality and Analysis :

- ✓ Name: The brand and model of the car.
- ✓ Location: The location in which the car is being sold or is available for purchase.
- ✓ Year: The year or edition of the model.
- ✓ Kilometers_Driven: The total kilometres driven in the car by the previous owner(s) in KM.
- ✓ Fuel_Type: The type of fuel used by the car.
- ✓ Transmission: The type of transmission used by the car.
- ✓ Owner_Type: Whether the ownership is Firsthand, Second hand or other.
- ✓ Mileage: The standard mileage offered by the car company in kmpl or km/kg.
- ✓ Engine: The displacement volume of the engine in cc.
- ✓ Power: The maximum power of the engine in bhp.
- ✓ Seats: The number of seats in the car.
- ✓ New_Price: The price of a new car of the same model.
- ✓ Price: The price of the used car in INR Lakhs.

Dataset number of columns and rows.

```
In [149]: car_train.shape
```

```
Out[149]: (5975, 17)
```

```
In [150]: car_test.shape
```

```
Out[150]: (1201, 12)
```

Top 5 values from the test dataset.

```
In [7]: car_test.head()
```

```
Out[7]:
```

	Unnamed: 0	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	Engine	Power	Seats	New_Price
0	0	Maruti Alto K10 LXI CNG	Delhi	2014	40929	CNG	Manual	First	32.26 km/kg	998 CC	58.2 bhp	4.0	NaN
1	1	Maruti Alto 800 2016-2019 LXI	Coimbatore	2013	54493	Petrol	Manual	Second	24.7 kmpl	796 CC	47.3 bhp	5.0	NaN
2	2	Toyota Innova Crysta Touring Sport 2.4 MT	Mumbai	2017	34000	Diesel	Manual	First	13.68 kmpl	2393 CC	147.8 bhp	7.0	25.27 Lakh
3	3	Toyota Etios Liva GD	Hyderabad	2012	139000	Diesel	Manual	First	23.59 kmpl	1364 CC	null bhp	5.0	NaN
4	4	Hyundai i20 Magna	Mumbai	2014	29000	Petrol	Manual	First	18.5 kmpl	1197 CC	82.85 bhp	5.0	NaN

Concise summary of the dataframe, data types, null value counts and column names.

```
In [8]: car_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6019 entries, 0 to 6018
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            6019 non-null  int64
1   Name                  6019 non-null  object
2   Location              6019 non-null  object
3   Year                  6019 non-null  int64
4   Kilometers_Driven     6019 non-null  int64
5   Fuel_Type             6019 non-null  object
6   Transmission          6019 non-null  object
7   Owner_Type            6019 non-null  object
8   Mileage               6017 non-null  object
9   Engine               5983 non-null  object
10  Power                 5983 non-null  object
11  Seats                 5977 non-null  float64
12  New_Price             824 non-null   object
13  Price                 6019 non-null  float64
dtypes: float64(2), int64(3), object(9)
memory usage: 658.5+ KB
```

Feature variables analysis:

```
In [10]: car_test.Mileage.min()
Out[10]: '0.0 kmpl'

In [11]: car_test.Fuel_Type.value_counts()
Out[11]: Diesel      647
         Petrol     579
         CNG         6
         LPG         2
         Name: Fuel_Type, dtype: int64
```

To view some basic statistical details like percentile, mean, std etc. of a data frame

```
In [17]: car_train.describe()
Out[17]:
```

	Unnamed: 0	Year	Kilometers_Driven	Seats	Price
count	6019.000000	6019.000000	6.019000e+03	5977.000000	6019.000000
mean	3009.000000	2013.358199	5.873838e+04	5.278735	9.479468
std	1737.679967	3.269742	9.126884e+04	0.808840	11.187917
min	0.000000	1998.000000	1.710000e+02	0.000000	0.440000
25%	1504.500000	2011.000000	3.400000e+04	5.000000	3.500000
50%	3009.000000	2014.000000	5.300000e+04	5.000000	5.640000
75%	4513.500000	2016.000000	7.300000e+04	5.000000	9.950000
max	6018.000000	2019.000000	6.500000e+06	10.000000	160.000000

Dropping column new_price from both test and train dataset.

```
In [12]: #The new price column is missing in both train and test set, so we can remove it
car_train=car_train.drop('New_Price',axis=1)
car_test=car_test.drop('New_Price',axis=1)
```

Dropping NA values

```
In [23]: car_train=car_train.dropna()
```

```
In [24]: car_train.shape
```

```
Out[24]: (5975, 13)
```

The pairwise correlation of all columns in the dataframe

```
In [19]: car_train.corr()
```

Out[19]:

	Unnamed: 0	Year	Kilometers_Driven	Seats	Price
Unnamed: 0	1.000000	0.002354	-0.008734	-0.010832	-0.020275
Year	0.002354	1.000000	-0.173048	0.012333	0.305327
Kilometers_Driven	-0.008734	-0.173048	1.000000	0.083113	-0.011493
Seats	-0.010832	0.012333	0.083113	1.000000	0.052225
Price	-0.020275	0.305327	-0.011493	0.052225	1.000000

```
In [20]: car_test.corr()
```

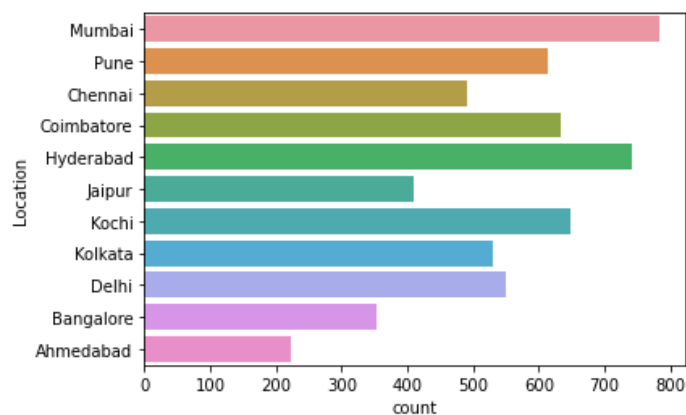
Out[20]:

	Unnamed: 0	Year	Kilometers_Driven	Seats
Unnamed: 0	1.000000	0.027562	-0.043032	-0.034890
Year	0.027562	1.000000	-0.455227	-0.011982
Kilometers_Driven	-0.043032	-0.455227	1.000000	0.219258
Seats	-0.034890	-0.011982	0.219258	1.000000

Location count plot.

```
In [26]: import seaborn as sns
sns.countplot(y='Location',data=car_train)
print(car_train.Location.value_counts(normalize=True)*100)
```

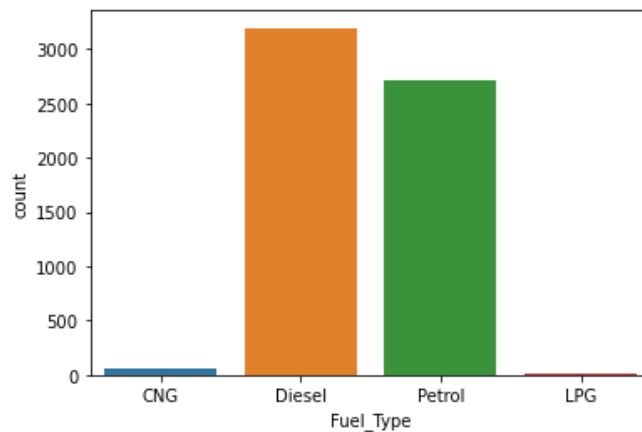
```
Mumbai      13.121339
Hyderabad   12.401674
Kochi       10.845188
Coimbatore  10.610879
Pune        10.259414
Delhi       9.188285
Kolkata     8.870293
Chennai     8.200837
Jaipur      6.861925
Bangalore   5.907950
Ahmedabad   3.732218
Name: Location, dtype: float64
```



Fuel type count plot

```
In [28]: sns.countplot(x='Fuel_Type',data=car_train)  
print(car_train.Fuel_Type.value_counts())
```

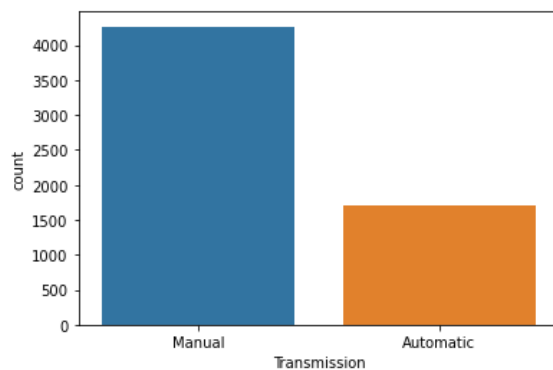
```
Diesel    3195  
Petrol    2714  
CNG        56  
LPG        10  
Name: Fuel_Type, dtype: int64
```



Transmission count plot.

```
In [29]: sns.countplot(x='Transmission',data=car_train)  
print(car_train.Transmission.value_counts(normalize=True)*100)
```

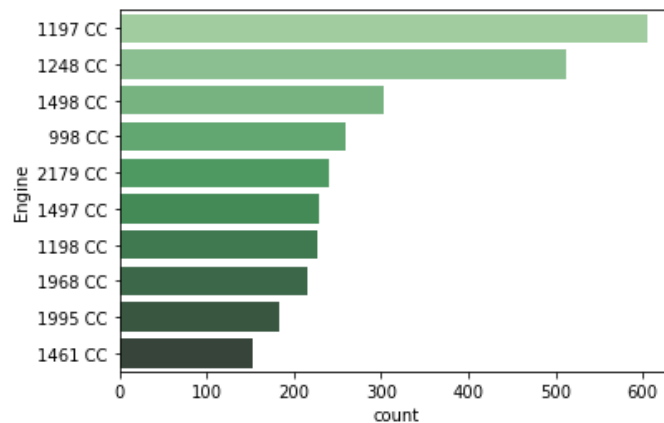
```
Manual    71.39749  
Automatic 28.60251  
Name: Transmission, dtype: float64
```



Engine count plot.

```
In [32]: sns.countplot(y="Engine", data=car_train, palette="Greens_d",  
                    order=car_train.Engine.value_counts().iloc[:10].index)
```

```
Out[32]: <AxesSubplot:xlabel='count', ylabel='Engine'>
```



Cleaning the values in the following columns

Power: 58.16 bhp -> 58.16

Milage: 19.67 kmpl -> 19.67

Engine: 998 CC -> 998

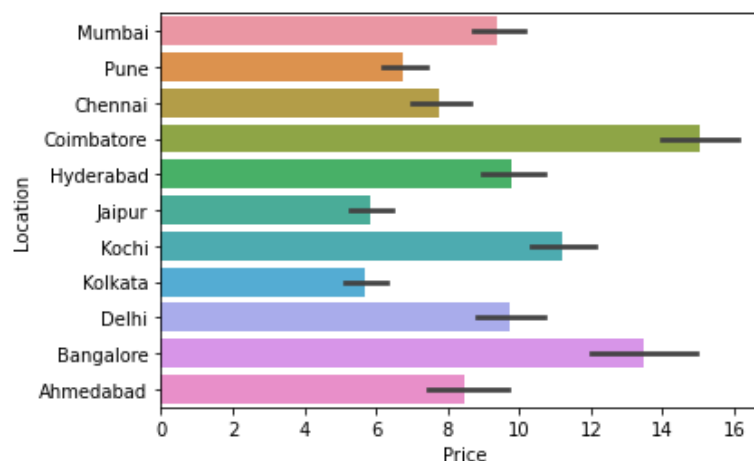
```
In [39]: car_train['power_n']=car_train.Power.str.extract(r'(\d+\.\d+)').astype('float')  
car_train['milage_n']=car_train.Mileage.str.extract(r'(\d+\.\d+)').astype('float')  
car_train['Engine_n']=car_train.Engine.str.extract(r'(\d+\.\d+)').astype('int')  
car_train['seat_n']=car_train.Seats.astype('int')
```

Bivariate Analysis

Price vs Location

```
sns.barplot(y="Location", x="Price", data=car_train)  
  
#here we are seeing the median price at each location
```

```
Out[50]: <AxesSubplot:xlabel='Price', ylabel='Location'>
```

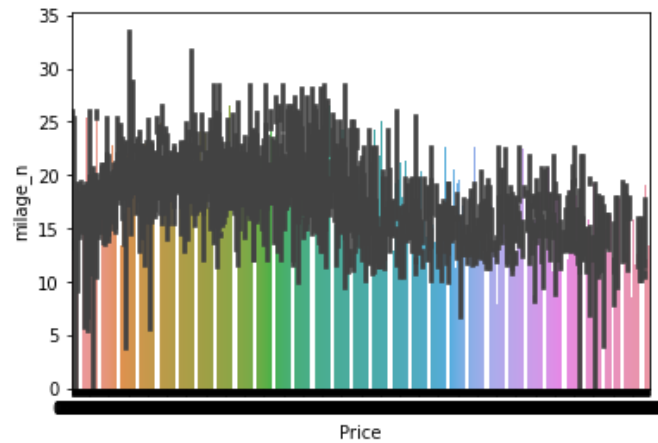


Milage vs Price

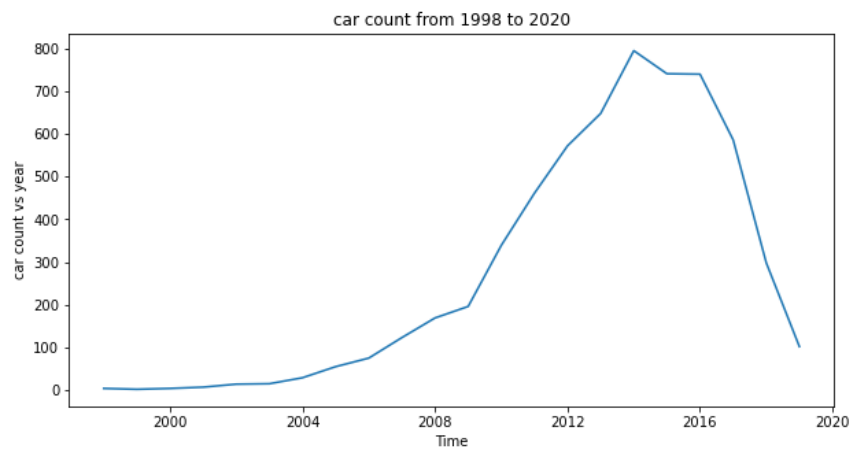
In [51]: *#Milage vs price*

```
sns.barplot(y="milage_n", x="Price", data=car_train)
```

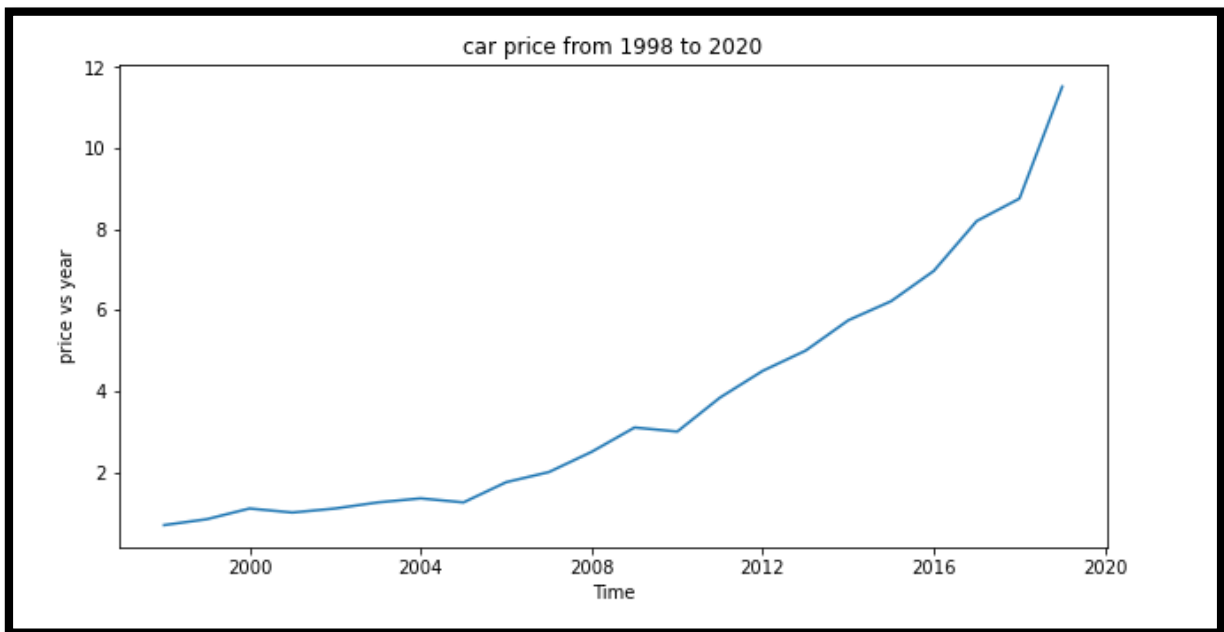
Out[51]: <AxesSubplot:xlabel='Price', ylabel='milage_n'>



Count of car vs Years

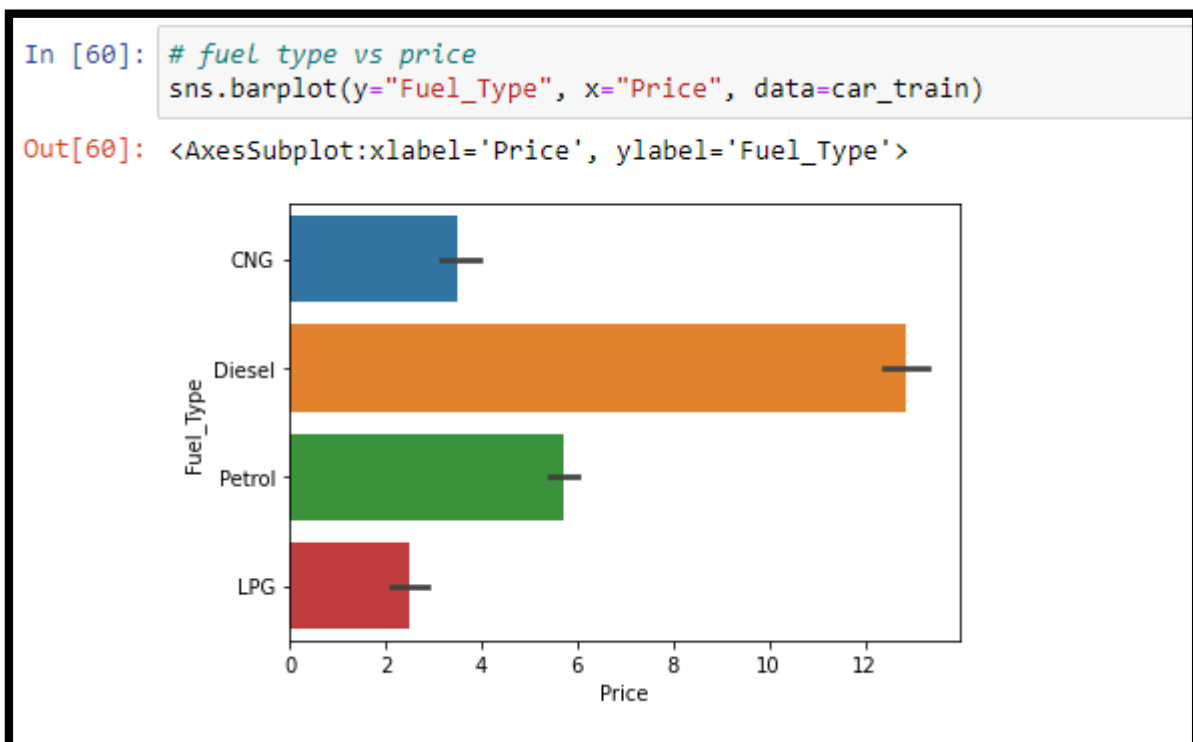


Car Price vs Years



It is been clear from the graph over the period of time the price has increased.

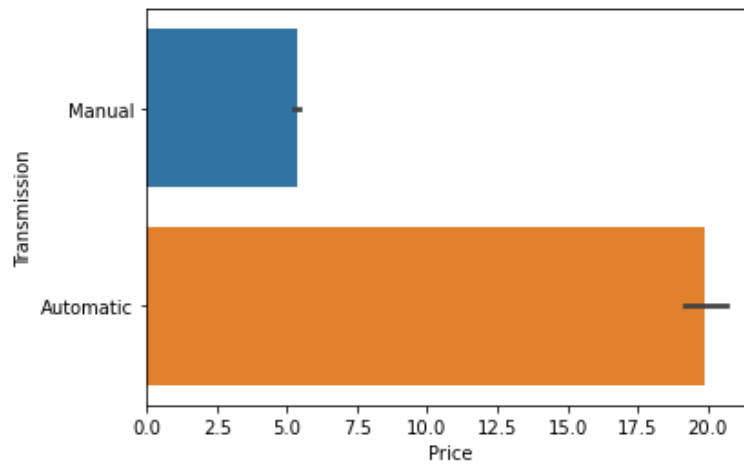
Fuel Type vs Price



Transmission vs Price

```
In [61]: #Transmission vs Price  
sns.barplot(y="Transmission", x="Price", data=car_train)
```

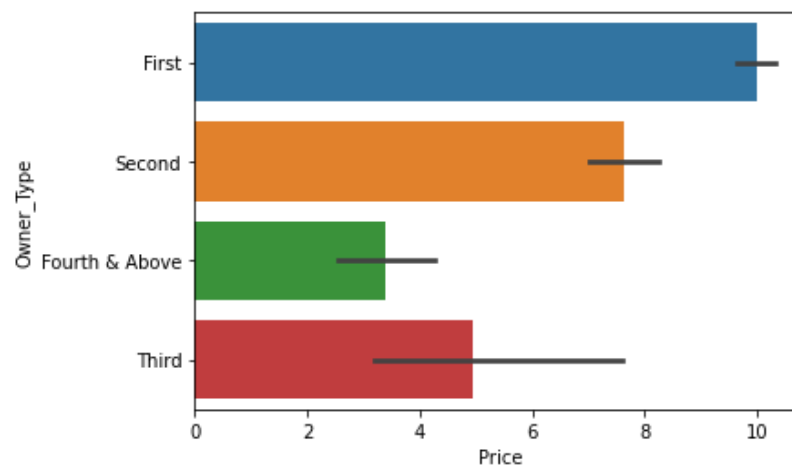
```
Out[61]: <AxesSubplot:xlabel='Price', ylabel='Transmission'>
```



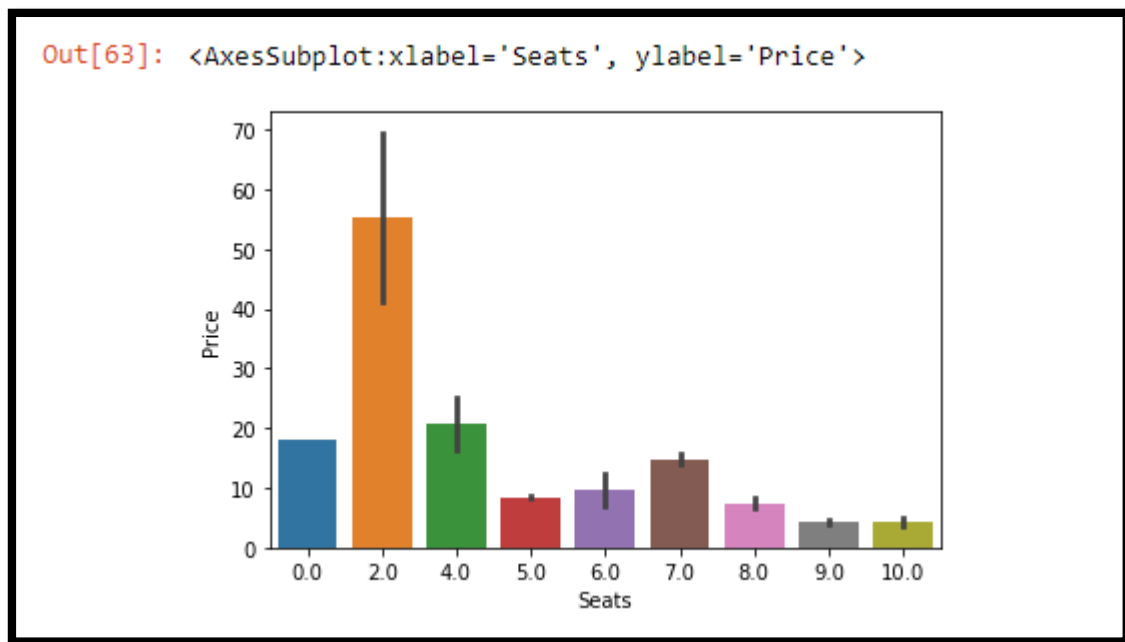
Owner Type vs Price

```
In [62]: #Owner_Type vs price  
sns.barplot(y="Owner_Type", x="Price", data=car_train)
```

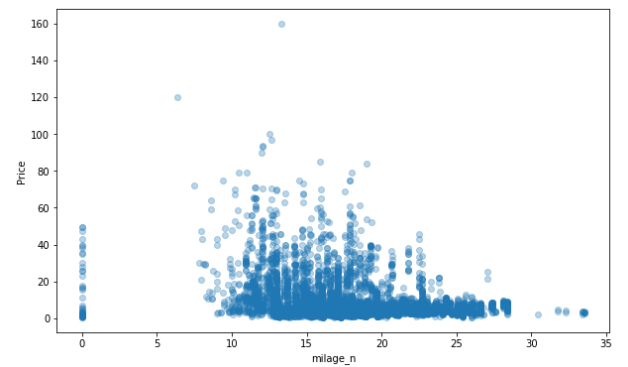
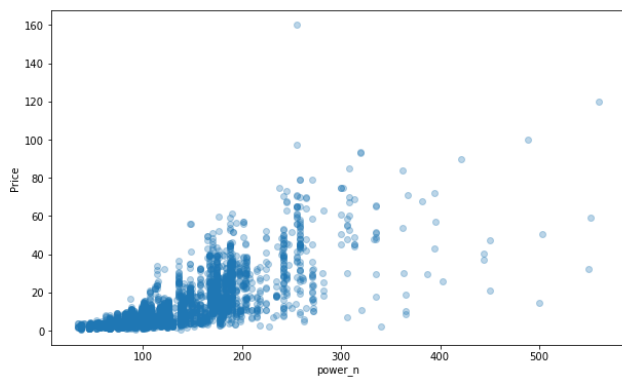
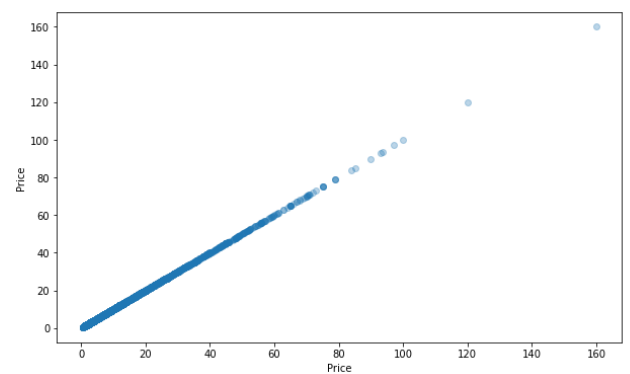
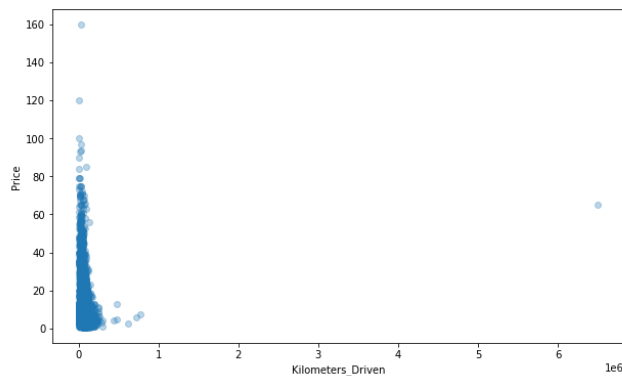
```
Out[62]: <AxesSubplot:xlabel='Price', ylabel='Owner_Type'>
```



No. of seats in the car Vs Price



Outliers



Removing outliers using BaseEstimator

```
In [69]: from sklearn.base import BaseEstimator,TransformerMixin,RegressorMixin
```

```
In [70]: class RemoveOutliers(BaseEstimator,TransformerMixin):
        """This class removes outliers from data.
        """

        def fit (self,X,y=None):
            return self

        def transform(self,X,y=None):
            X=X[X['Kilometers_Driven']< 262000]

            X=X[X['Price']<=100]

            X=X[X['power_n']<= 530]

            X=X[X['Engine_n']<= 5900 ]

            return X
```

```
In [71]: data1=RemoveOutliers().fit_transform(car_train)
```

Feature Transformation

```
In [86]: class FeaturesTransformer(BaseEstimator,TransformerMixin):
        """This class trnsforms numerical featureess in the dataset.
        |transformations are hard coded.
        """

        def fit(self,X,y=None):
            return self

        def transform(self,X,y=None):
            import numpy as np
            from scipy.special import boxcox1p
            X['Kilometers_Driven']=X['Kilometers_Driven'].apply(lambda x: boxcox1p(x,0.33))

            X['power_n']=X['power_n'].apply(lambda x:np.log(x) )

            X['Engine_n']=X['Engine_n'].apply(lambda x: np.log(x))

            X['milage_n']=X['milage_n'].apply(lambda x: np.log1p(x)**4)

            return X
```

Creating dummy variables

Dummy Variables act as indicators of the presence or absence of a category in a Categorical Variable. The usual convention dictates that 0 represents absence while 1 represents presence. The conversion of Categorical Variables into Dummy Variables leads to the formation of the two-dimensional binary matrix where each column represents a particular category.

One Hot Encoding.

A one hot encoding is a representation of categorical variables as binary vectors.

This first requires that the categorical values be mapped to integer values.

Then, each integer value is represented as a binary vector that is all zero values except the index of the integer, which is marked with a 1.

```
In [98]: #as Transmission is an nominal variable lets perform onehotencoding
Transmission = data3[["Transmission"]]

Transmission = pd.get_dummies(Transmission,drop_first=True)

Transmission.head()

# Location
location = data3[["Location"]]

location = pd.get_dummies(location,drop_first=True)

location.head()

#fuel_type
Fuel_Type = data3[["Fuel_Type"]]

Fuel_Type = pd.get_dummies(Fuel_Type,drop_first=True)

Fuel_Type.head()
```

Out[98]:

	Fuel_Type_Diesel	Fuel_Type_LPG	Fuel_Type_Petrol
0	0	0	0
1	1	0	0
2	0	0	1
3	1	0	0
4	1	0	0

Concatenate dataframe >data3+ Location + Transmission + Fuel_Type

```

In [99]: #Concatenate dataframe >data3+ company + Location + Transmission + Fuel_Type
data_train = pd.concat([data3 ,location ,Transmission,Fuel_Type ], axis = 1)

In [100]: data_train.head(3)
Out[100]:

```

	Unnamed: 0	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Price	power_n	milage_n	...	Location_Hyderabad	Location_Jaipur
0	0	Mumbai	2010	118.422565	CNG	Manual	1	1.75	4.063198	121.173895	...	0	0
1	1	Pune	2015	97.827204	Diesel	Manual	1	12.50	4.837868	84.142516	...	0	0
2	2	Chennai	2011	101.730606	Petrol	Manual	1	4.50	4.485260	76.239001	...	0	0

3 rows x 26 columns

III. Features Pre-Processing :

Standardization of a dataset is a common requirement for many machine learning estimators: they might behave badly if the individual features do not more or less look like standard normally distributed data

```

In [103]: features_c
Out[103]: ['Kilometers_Driven', 'power_n', 'milage_n', 'Engine_n']

In [104]: scaled_features = data_train.copy()

In [105]: from sklearn.preprocessing import StandardScaler

In [106]: col_names = ['Kilometers_Driven','power_n','milage_n','Engine_n']
features = scaled_features[col_names]
scaler = StandardScaler().fit(features.values)
features = scaler.transform(features.values)

In [107]: scaled_features[col_names] = features
#print(scaled_features)

In [108]: scaled_features.head(2)
Out[108]:

```

	Unnamed: 0	Year	Kilometers_Driven	Owner_Type	Price	power_n	milage_n	Engine_n	seat_n	Location_Bangalore	...	Location_Hyderabad	Location_Jai
0	0	2010	0.641616	1	1.75	-1.541884	2.034639	-1.354647	5	0	...	0	
1	1	2015	-0.288928	1	12.50	0.323319	0.399417	-0.023337	5	0	...	0	

2 rows x 23 columns

IV. Format of the Dataset : CSV

B. Identification of Learning Model (Supervised Learning)

I. Algorithm used : RandomForest

The random forest is a classification algorithm consisting of many decisions trees. It uses **bagging and feature randomness when building each individual tree to try to create an uncorrelated forest of trees** whose prediction by committee is more accurate than that of any individual tree.

II. Methodology used :

One Hot Encoding.

A one hot encoding is a representation of categorical variables as binary vectors. This first requires that the categorical values be mapped to integer values. Then, each integer value is represented as a binary vector that is all zero values except the index of the integer, which is marked with a 1.

Creating dummy variables

Dummy Variables act as indicators of the presence or absence of a category in a Categorical Variable. The usual convention dictates that 0 represents absence while 1 represents presence. The conversion of Categorical Variables into Dummy Variables leads to the formation of the two-dimensional binary matrix where each column represents a particular category.

III. Model building, Training & Testing :

Model Training

```
In [122]: y=scaled_features['Price']  
X=scaled_features.drop('Price',axis=1)  
  
In [123]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 25)
```

Linear Regression

Linear Regression

```
In [127]: from sklearn.linear_model import LinearRegression
linear_reg = LinearRegression()
linear_reg.fit(X_train, y_train)
y_pred = linear_reg.predict(X_test)
print("Accuracy on Traing set: ",linear_reg.score(X_train,y_train)*100,"%")
print("Accuracy on Testing set: ",linear_reg.score(X_test,y_test)*100,"%")

Accuracy on Traing set: 66.48679499628179 %
Accuracy on Testing set: 67.8906639583188 %
```

Random Forest

RandomForest

```
In [128]: from sklearn.ensemble import RandomForestRegressor
reg_rf = RandomForestRegressor()
reg_rf.fit(X_train, y_train)
y_pred = reg_rf.predict(X_test)
print("Accuracy on Traing set: ",reg_rf.score(X_train,y_train)*100,"%")
print("Accuracy on Testing set: ",reg_rf.score(X_test,y_test)*100,"%")

Accuracy on Traing set: 98.58226128836158 %
Accuracy on Testing set: 86.73912912305634 %
```

IV. Model Accuracy, Prediction & Precession :

```
In [129]: y_pred = reg_rf.predict(X_test)

In [136]: from sklearn import metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error

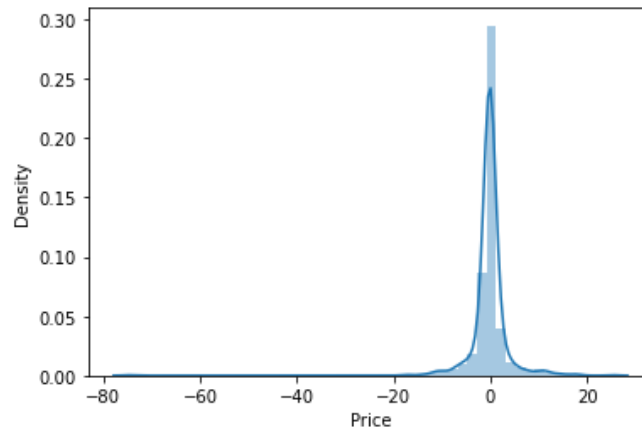
print("\t\tError Table")
print('Mean Absolute Error      : ', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error      : ', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error : ', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R Squared Error          : ', metrics.r2_score(y_test, y_pred)*100, '%')

Error Table
Mean Absolute Error      : 1.8073331632653065
Mean Squared Error      : 16.781610776030604
Root Mean Squared Error : 4.09653643655596
R Squared Error          : 86.73912912305634 %
```

```
In [131]: sns.distplot(y_test-y_pred)
plt.show()
```

C:\Users\Hp\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is deprecated. It will be removed in a future version. Please adapt your code to use either `display` or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



```
In [132]: from sklearn import metrics
y_pred = reg_rf.predict(X_test)
print('Accuracy Score:')
print(metrics.r2_score(y_test, y_pred)*100, "%")
```

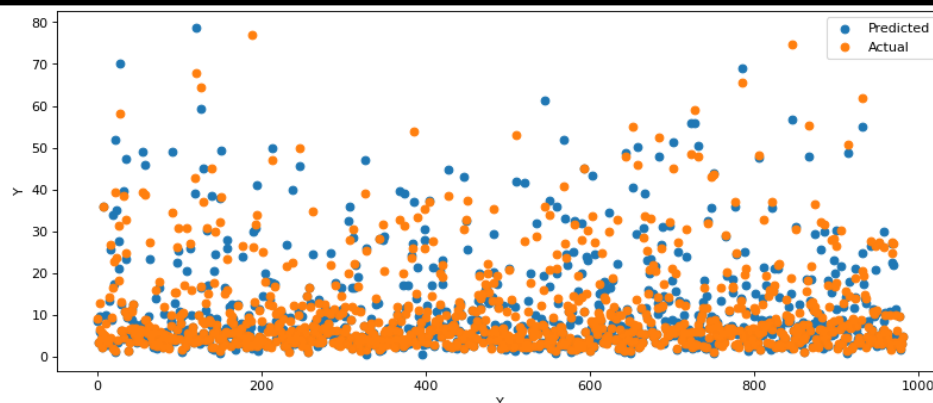
Accuracy Score:
86.73912912305634 %

```
In [133]: print('MAE:', metrics.mean_absolute_error(y_test, y_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R Squared Error : ', metrics.r2_score(y_test, y_pred)*100, "%")

print("Accuracy on Traing set: ", reg_rf.score(X_train, y_train)*100, "%")
```

MAE: 1.8073331632653065
MSE: 16.781610776030604
RMSE: 4.09653643655596
R Squared Error : 86.73912912305634 %
Accuracy on Traing set: 98.58226128836158 %

Predicted value vs Actual Value



C. Key Learning Outcomes :

After the completing this assignment I came to know that the data cleaning and standardization plays an important role for the accuracy of the model. Data cleaning, filling null, data standardization of values can abruptly increase the accuracy of the model. For the above case study, I had designed two models, one without standardization of the features, another one with standardisation of features and converting the objects to categorical variables using one hot encoding, the second model proved to be more accurate and precise as compared to the one without standardization and categorical conversion.

Viva Questions:

Random Forest and Linear Regression Accuracy

In the dataset we have 4-5 features which are majorly effecting the price of the car, also dataset contains features some of which are Categorical Variables and some of the others are continuous variable. Decision Tree is better than **Linear Regression**, since Trees can accurately divide the data based on Categorical Variables. Decision Trees are great for obtaining non-linear relationships between input features and the target variable.

Random Forest can handle messier data and messier relationships better than regression models. The averaging of the predicted values makes Random Forest better than a single Decision Tree hence improves its accuracy and reduces overfitting.

Random Forest and K-Means Clustering (Assigning value of k in Random Forest?)

The random forest algorithm is a supervised learning model; it uses labeled data to “learn” how to classify unlabeled data. This is the opposite of the K-means Cluster algorithm, it is an unsupervised learning model. The Random Forest Algorithm is used to solve both regression and classification problems, making it a diverse model that is widely used by engineers.

The k-means clustering algorithm assigns data points to categories, or clusters, by finding the mean distance between data points. It then iterates through this technique in order to perform more accurate classifications over time. Since you must first start by classifying your data into k categories.

Submitted By:

Aditya Sharma

AU18B1009